

An Empirical Analysis of Optimization for Max-Margin NLP

Jonathan K. Kummerfeld, Taylor Berg-Kirkpatrick, and Dan Klein

{jkk,tberg,klein}@cs.berkeley.edu

Some margin optimizers are better than others... Use this one!

There is substantial variation in the effectiveness of optimization methods for structured max-margin objectives. We investigated the behavior of a range of optimizers (dual and primal margin, likelihood, perceptron, and MIRA) for training high-performance systems for NER, coreference, parsing and summarization.

Symbol Definitions

w	Weights
q	Sum of squared subgradients
u	Last update for each weight
n	Number of updates
g	Subgradient
f	Feature index
η	Learning rate
δ	Small value to prevent $q[f] = 0$

```
function optimize-margin(data, iters):
```

```
    w[] = 0
    q[] = delta
    u[] = 0
    n = 0
    for iter in [1, iters]:
        for (x*, y*) in data:
            y' = argmax (score(x*, y) + L(y, y*))
                  (over y in Y(x*))
            g = features(y*) - features(y')
            q += g^2
            n += 1
            for f in (nonzero features in g):
                w[f] = update-active(w[f], g[f], q[f])
                w[f] = n
```

Main Loop

Loss augmented decoding, with tracking of the sum of squared subgradients. A range of methods can be implemented by varying L , update-active, and score. For example, the perceptron is: $L = 0$, score = $f \cdot w$, and update-active = $w + g$

AdaGrad Update

The AdaGrad update with L1 regularization. See the paper for L2.

```
function update-active(w, g, q):
```

```
    d = |w - g*eta / sqrt(q)| - C*eta / sqrt(q)
    return sign(w - g*eta / sqrt(q)) * max(0, d)
```

```
function update-catchup(w, q, t):
```

```
    return sign(w)*max(0, |w| - eta*C*t / sqrt(q))
```

```
function score(x, y):
```

```
    s = 0
    for f in features(x, y):
        w[f] = update-catchup(w[f], q[f], n - u[f])
        u[f] = n
        s += w[f]
    return s
```

Sparse updates

AdaGrad modifies every weight on every update. A fast way to implement this is to only update weights for nonzero features at first (in the main loop, above), and update other weights the next time they are needed(here).

Observations

- Max-margin generally outperformed MIRA, perceptron, likelihood.
 - Cutting plane methods learned more slowly.
 - OPS (Online Primal Subgradient) was easy, robust, and effective.
- Decoding dominated time for each iteration on most tasks. Sparse updates were crucial for OPS on some tasks.

Code Release

See nlp.cs.berkeley.edu for our learning library!

